

Part 6: Simple Linear Regression Theory

This lecture provides an overview of the most basic regression model, the **simple linear regression model**. There is a single **covariate/predictor** x used to predict the **response** Y . This model will serve as a building block for the more complex (and more useful) models we will present afterwards.

The simple linear regression model is typically written using notation

$$Y_i = \beta_0 + \beta_1 x_i + \epsilon_i, \quad i = 1, 2, \dots, n$$

where β_0 is the **intercept** and β_1 is the **slope** for the regression line.




Figure 1: Synthetic scatter plot with $n = 50$ and a fitted regression line, $Y = 3.15 + 1.65x$.

The ϵ_i are random variables capturing the **irreducible error**, the scatter around the regression line.

The exact assumptions placed on the ϵ_i will vary depending on the situation, but the base assumptions that are almost always associated with simple linear regression are as follows:

1. $E(\epsilon_i) = 0$ for all i

2. $V(\epsilon_i) = \sigma^2$ for all i

3. The ϵ_i are uncorrelated, i.e., $\text{Cov}(\epsilon_i, \epsilon_j) = 0$ for all $i \neq j$.

Additional assumptions, **if justified**, lead to stronger results. In particular, one often assumes that the ϵ_i are normally distributed.

Example: This model is sometimes written as

$$Y = \alpha + \beta x + \epsilon.$$

This notation is the origin for references to the “beta” for a stock.

Starbucks Corporation (NASDAQ:SBUX)

56.18	-0.13 (-0.23%)	Range	55.99 - 56.65	Div/yield	0.20/1.42
Sep 2 - Close		52 week	52.63 - 64.00	EPS	1.79
NASDAQ real-time data - Disclaimer		Open	56.52	Shares	1.47B
Currency in USD		Vol / Avg.	7.44M/8.35M	Beta	0.81
		Mkt cap	82.32B	Inst. own	70%
		P/E	31.43		

Figure 2: Screen shot from Google Finance, taken September 3, 2016.

In particular, the “beta” is found by performing a simple linear regression where the response is the excess returns for that equity, and the predictor is the excess returns for the market (as measured by some standard index such as S & P 500).

- 1 Figure 3 shows an example of the use of simple linear regression with the
2 objective of calculating the “beta” for Google (GOOG), using daily market
3 from 2014.




Figure 3: Daily excess returns for GOOG versus daily excess returns for S & P 500 for 2014.

Estimation via a Sample

In practice, we rely on a sample of n pairs (x_i, y_i) for $i = 1, 2, \dots, n$, assumed to be drawn from a larger population. These will be used to estimate β_0, β_1 and σ^2 . The standard notation for these estimates will be $\hat{\beta}_0, \hat{\beta}_1$, and $\hat{\sigma}^2$.

The sample is naturally depicted in a **scatter plot**, with example shown in Figure 4.




Figure 4: A simple scatter plot, with $n = 20$.

1 Some additional notation we should establish immediately: For any esti-
2 mated regression line, the **fitted values** are

3
$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i, \quad i = 1, 2, \dots, n$$

4 and the **residuals** are

5
$$\hat{\epsilon}_i = y_i - \hat{y}_i, \quad i = 1, 2, \dots, n.$$

6 **Exercise:** Use the scatter plot below, with a fit regression line superim-
7 posed, to demonstrate the fitted values and the residuals.




Figure 5: A simple scatter plot, with $n = 20$, with a regression line fit shown.

The Sample Covariance and Correlation

★ **Exercise:** Contrast covariance and correlation.

Recall that the covariance (and a related quantity, correlation) are measures of the **linear** association between two random variables X and Y :

$$\text{Cov}(X, Y) = E[(X - E(X))(Y - E(Y))]$$

and

$$\text{Corr}(X, Y) = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y}$$

Not surprisingly, sample versions of these quantities are relevant to simple linear regression.

First, the **sample covariance** between observations of n pairs (x_i, y_i) :

$$s_{xy} = \left(\frac{1}{n-1} \right) \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$$

Note the similarity between this and the usual sample variance:

$$s_x^2 = \left(\frac{1}{n-1} \right) \sum_{i=1}^n (x_i - \bar{x})^2$$

Then, the **sample correlation** is

$$r_{xy} = \frac{s_{xy}}{s_x s_y}.$$

- 1 This quantity is typically just denoted with r . Figure 6 shows some exam-
 2 ples of scatter plots, and the value of r for each. It is useful for building
 3 some intuition as to what different values of r mean.





Figure 6: Examples of scatter plots, and corresponding values for r .

- 4 **Exercise:** Consider the scatter plot below. What, approximately, is the sam-
 5 ple correlation?



The Parameter Estimates

★ **Exercise:** What route would you use to derive $\hat{\beta}_0$ and $\hat{\beta}_1$?

If we were to follow through with this, we would find that the optimal estimates would result from minimizing the **residual sum of squares (RSS)**:

$$\text{RSS}(\beta_0, \beta_1) = \sum_{i=1}^n (y_i - (\beta_0 + \beta_1 x_i))^2$$

And a little bit of calculus would show that

$$\hat{\beta}_1 = \frac{\sum_i (y_i - \bar{y})(x_i - \bar{x})}{\sum_i (x_i - \bar{x})^2} = \frac{s_{xy}}{s_x^2} = r_{xy} \left(\frac{s_y}{s_x} \right)$$

and

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}.$$

1 ★ **Exercise:** Why do we use **least squares** to estimate the parameters in
2 regression?

3 _____

4 _____

5 _____

6 _____

7 _____

8 _____

9 _____

10 ★ **Exercise:** Suppose that the roles of the predictor and response variables
11 were interchanged in this regression model. What would be the effect on
12 the slope of the regression line?

13 _____

14 _____

15 _____

16 _____

17 _____

18 _____

19 _____

- Figure 7 illustrates this idea. On the same scatter plot, two regression lines are shown: The “dashed-dotted” line is the regression line for predicting GOOG excess return from the S&P 500 excess return. The “dashed” line is the regression line for predicting S&P 500 excess return from GOOG excess return. The solid line is the **SD Line**. This has a slope equal to s_y/s_x .




Figure 7: Daily excess returns for GOOG versus daily excess returns for S & P 500 for Jan. to Nov. 2012.

Estimating σ^2

The MLE for σ^2 can be derived to be

$$\hat{\sigma}_{\text{MLE}}^2 = \left(\frac{1}{n}\right) \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \text{RSS}/n.$$

This estimator is biased, but it can be adjusted to obtain an unbiased estimator

$$\hat{\sigma}^2 = \left(\frac{n}{n-2}\right) \sigma_{\text{MLE}}^2 = \left(\frac{1}{n-2}\right) \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \text{RSS}/(n-2).$$

This is the estimator we will typically use, and what is reported by R and other software packages.

Exercise: Is there intuition as to why we divide by $(n-2)$?




Figure 8: The geometric interpretation of least squares.

Figure 8 is a very useful geometric depiction of what happens when we do least squares. The vector of fitted values $\hat{y}$ is the orthogonal projection of y onto the space spanned by the vectors 1 and x . (The rectangle represents all of the possible vectors that can be formed from linear combinations of 1 and x .)

★ **Exercise:** What is the sum of the residuals in OLS?

Inference Regarding β_1

Within the context of this simple model, the primary objective is to perform inference regarding β_1 , since it captures information regarding the relationship (or lack thereof) between the response and the single predictor.

Here are some key results in this direction. Keep in mind that these are all true under the assumptions put forth for the simple linear regression model.

- $\hat{\beta}_1$ is an unbiased estimator of β_1 .

- The variance of $\hat{\beta}_1$ is

$$V(\hat{\beta}_1) = \frac{\sigma^2}{(n-1)s_x^2} \approx \frac{\hat{\sigma}^2}{(n-1)s_x^2}.$$

This means that the standard error is

$$\text{SE}(\hat{\beta}_1) \approx \frac{\hat{\sigma}}{s_x \sqrt{(n-1)}}.$$

- Suppose that the true value of the slope is β_{10} . Then

$$T = \frac{\hat{\beta}_1 - \beta_{10}}{\text{SE}(\hat{\beta}_1)}$$

has the t -distribution with $(n-2)$ degrees of freedom. Of course, when n is sufficiently large, T will be approximately standard normal.

1 **Exercise:** Construct a $100(1 - \alpha)\%$ confidence interval for β_1 .

2 _____

3 _____

4 _____

5 _____

6 _____

7 _____

8 _____

9 _____

10 _____

11 _____

12 _____

A common (overused) objective is to test for evidence if one can claim that $\beta_1 \neq 0$. It is important to be mindful of the following fact: If the confidence interval for β_1 includes zero, this is often not because $\beta_1 = 0$, but instead because the standard error for $\hat{\beta}_1$ is too large to detect the deviation from zero.

★ **Exercise:** How can the standard error of $\hat{\beta}_1$ be decreased?

One of the most common mistakes made when constructing models is to blindly remove a predictor from the model solely because there is not strong evidence that its coefficient does not equal zero. Model building (the process of deciding which covariates to include) should comprise many factors, including our knowledge of the situation, the importance that we are placing on constructing a **parsimonious** model, **and** the results of formal statistical inference procedures.

The Coefficient of Determination, R^2

When working with regression models it is very common to hear references to “the R^2 for the model.”

This is also called the **coefficient of determination**.

The idea is simple: R^2 equals the proportion of the sum of squared response which is accounted for by the model that has been fit, relative to the model with no covariates.

In particular:

$$R^2 = 1 - \left[\sum_{i=1}^n (y_i - \hat{y}_i)^2 \middle/ \sum_{i=1}^n (y_i - \bar{y})^2 \right]$$

Note that $0 \leq R^2 \leq 1$.

In the case of simple linear regression (a single predictor), R^2 does indeed equal the sample correlation (r) squared.

R^2 is overused. In particular, note that it is possible for R^2 to be close to one, but the model be a terrible fit. And it is possible for a great-fitting model to have R^2 close to zero. See Figure 9.




Figure 9: The scatter plot on the left has $R^2 \approx 0.05$, while that on the right has $R^2 \approx 0.95$. The data in the left plot were simulated from the normal linear model.

In short, R^2 should not be relied upon as a diagnostic to judge the quality of the model fit. **If** the model is a good fit, **then** R^2 says something about the **predictive power** of the model.

★ **Exercise:** Suppose that, by mistake, each observed pair (response and predictor) is included twice in the analysis. What is the effect on $\hat{\beta}_i$, the t -statistics, and R^2 ?