

Part 9: Power Transformations

In this brief part we will consider families of transforming a sample with the explicit objective of maximizing the quality of the fit of a normal distribution to that sample.

The motivations for such a transformation include (1) to make a skewed distribution more symmetric and “nice”; (2) to allow for the use of the normal model in a case where it otherwise may not be clear how to model the variable; and (3) to transform the response of a regression model so that the normal approximation is a better fit for the model error.

Box-Cox Transformation

This is **only** useful in the case where the response is **strictly positive**.

Define Box-Cox **power transformation** function:

$$y^{(\lambda)} = \begin{cases} \frac{y^\lambda - 1}{\lambda} & \text{if } \lambda \neq 0 \\ \log(y) & \text{if } \lambda = 0 \end{cases}$$

This defines a wide range of possible transformations for the response, including $\sqrt{y}$, y^2 , $\log(y)$, $1/\sqrt{y}$, and so forth.

Exercise: Why is the case $\lambda \neq 0$ defined in the manner it is, instead of just as y^λ ?

In Python, this is implemented via the function `boxcox()`, which is part of the package `scipy.stats`.

```
In [1]: from scipy import stats
        import matplotlib.pyplot as plt
        import numpy as np
```

Let's start with a synthetic example.


```
In [2]: x = np.random.gamma(1, 4, 100)
```

We note that this is far from a normal shape.

```
In [3]: from statsmodels.graphics.gofplots import qqplot

fig, axs = plt.subplots(1, 2, figsize=[10, 4.5])
axs[0].hist(x, density=True)
axs[0].set_xlabel("x")
axs[0].set_ylabel("Density")

qqplot(x, line="s", ax=axs[1])
plt.show()
```




Now we call the `boxcox` function. By specifying `lambda=None`, we ask the procedure to find the optimal value of λ via maximum likelihood. Specifying `alpha=0.05` means that the procedure will return a 95% confidence interval for λ .

```
In [4]: xt, lam, CI = stats.boxcox(x,  
                                     lambda=None, alpha=0.05)  
  
        print(lam)  
        print(CI)  
  
0.12639088708173513  
(-0.02035547883356052, 0.2772495624156657)
```

```
In [5]: fig, axs = plt.subplots(1, 2, figsize=[10,4.5])
        axs[0].hist(xt,density=True)
        axs[0].set_xlabel("x, transformed")
        axs[0].set_ylabel("Density")

        qqplot(xt,line="s",ax=axs[1])
        plt.show()
```



Comment: It is standard to choose a “nice” value of λ from the 95% confidence interval and use that for the transformation. For example, if 0.5 is in the confidence interval, one would often prefer to use that instead of $\lambda = 0.43$, because the square root transform is easier to interpret. Likewise, if zero is in the interval, the log transform is typically utilized.

Exercise: Consider the volume data below.


```
In [6]: import yfinance as yf

        SBUX = yf.download("SBUX", start="2019-01-01",
                           end="2019-08-31")

        fig, axs = plt.subplots(1, 2, figsize=[10,4.5])
        axs[0].hist(SBUX['Volume'],density=True)
        axs[0].set_xlabel("SBUX Volume")
        axs[0].set_ylabel("Density")

        qqplot(SBUX['Volume'],line="s", ax=axs[1])
        plt.show()
```

[*****100%*****] 1 of 1 down



Try the Box-Cox transformation in this case. What happens?

Box-Cox for Regression

The original motivation for this type of transformation was in the context of the linear regression model

$$Y^{(\lambda)} = \beta_0 + \sum_{j=1}^p \beta_j x_j + \epsilon$$

The question posed by Box and Cox was as follows: What choice of λ yields the best fitting model, if we want to assume the errors are i.i.d. Normal?

Through a maximum likelihood procedure, the best λ can be found, and then this choice is used to transform the response prior to fitting the full regression model.

Exercise: The Python function `boxcox()` described above operates on a single vector. How could that be used in the context of this regression problem?

Calling R from Python

For all of the appeal of Python, R is definitely better suited for certain statistical and machine learning procedures. It is possible to take advantage of these capabilities directly through Python, using the package `rpy2`.

There are different ways of doing this. Here we will show an approach which allows you to embed chunks of R code within the Jupyter notebook, and pass objects back and forth as needed.

First, load the needed package, and load the interactive extension using the following two commands.

```
In [7]: import rpy2
        %load_ext rpy2.ipython
```

R does not like variable names with spaces in them.


```
In [8]: SBUX = SBUX.rename(columns=\
                           {"Adj Close": "AdjClose"})
```

The R function `boxcox` that is part of the package `MASS` implements the Box-Cox procedure. Note that here, we only need to look at the plot, so no objects are returned from the R code. To return objects, use the argument `-o` in the first line. We see examples of that later.

```
In [9]: %%R -i SBUX
```

```
library(MASS)
```

```
boxcox (Volume~AdjClose, data=SBUX)
```



Exercise: What conclusion do you draw from this analysis?

Exercise: Fit the linear model using an appropriate transformation of the response. Assess the results.

The Yeo-Johnson Transformation

The Box-Cox Transformation suffers from the limitation that it can only be applied to strictly positive variables. A generalization that works for any variable is the **Yeo-Johnson Transformation**.

The Yeo-Johnson Power Transformation is

$$y^{(\lambda, \lambda)} = \begin{cases} ((y + 1)^\lambda - 1) / \lambda, & \text{if } \lambda \neq 0, y \geq 0 \\ \log(y + 1), & \text{if } \lambda = 0, y \geq 0 \\ - \left[(-y + 1)^{2-\lambda} - 1 \right] / (2 - \lambda), & \text{if } \lambda \neq 2, y < 0 \\ - \log(-y + 1), & \text{if } \lambda = 2, y < 0 \end{cases}$$

It appears somewhat messy, but the intention is much like for the Box-Cox Transformation: As λ varies, we get a wide range of power and log transformations of the data; the optimal transformation can be sought via maximum likelihood.

In Python this is implemented in `scipy.stats.yeojohnson`.

In []: