

Lecture 4: Regularized Regression - Lasso

Max G'Sell
Machine Learning I: 46-926

November 8, 2020

Announcements

- ▶ First quiz was turned in yesterday. How did that work out?

Today: Finishing basics, Linear Regression and Regularization

- ▶ Extending linear regression to work in high dimensions
 - ▶ Lasso
 - ▶ Model selection, cross-validation (unlikely)

Review: Supervised learning

Observe $(X_1, Y_1), \dots, (X_n, Y_n) \stackrel{iid}{\sim} \mathcal{P}$. Estimate $\hat{\mu}(x)$ so that $\mathbb{E}\mathcal{L}(Y, \hat{\mu}(X))$ is small for new $(X, Y) \sim \mathcal{P}$.

If we are minimizing expected squared prediction error:

- If we actually knew $\mathcal{P}$, the best we could do is

$$\mu(x) = \mathbb{E}(Y|X = x)$$

- If we are estimating $\hat{\mu}$, then

$$\mathbb{E}(\mathcal{L}(Y, \hat{\mu}(X)) \mid X = x) = \sigma_x^2 + \text{bias}(\hat{\mu}(x))^2 + \text{var}(\hat{\mu}(x))$$

Review: Linear regression


We can think of linear regression as one way to approximate $\mathbb{E}(Y|X = x)$ across x values, imagining:

$$\mathbb{E}(Y|X = x) \approx x^T \beta$$

This can work well if the linear model is not a bad approximation.

However, we noticed that, as the number of variables, p , grows large, the variance of linear regression grows rapidly. This leads to poor predictions **even with the linear model is true!**

Linear regression in high dimensions



This is a simulation from a true linear model with independent Gaussian errors – the dream! (Also the plot from your homework.)

Ridge regression

Ridge regression is like least squares but shrinks the estimated coefficients towards zero. Given a response vector $y \in \mathbb{R}^n$ and a predictor matrix $X \in \mathbb{R}^{n \times p}$, the ridge regression coefficients are defined as

$$\begin{aligned}\hat{\beta}^{\text{ridge}} &= \operatorname{argmin}_{\beta \in \mathbb{R}^p} \sum_{i=1}^n (y_i - x_i^T \beta)^2 + \lambda \sum_{j=1}^p \beta_j^2 \\ &= \operatorname{argmin}_{\beta \in \mathbb{R}^p} \underbrace{\|y - X\beta\|_2^2}_{\text{Loss}} + \lambda \underbrace{\|\beta\|_2^2}_{\text{Penalty}}\end{aligned}$$

Ridge regression


$$\begin{aligned}\hat{\beta}^{\text{ridge}} &= \operatorname{argmin}_{\beta \in \mathbb{R}^p} \sum_{i=1}^n (y_i - x_i^T \beta)^2 + \lambda \sum_{j=1}^p \beta_j^2 \\ &= \operatorname{argmin}_{\beta \in \mathbb{R}^p} \underbrace{\|y - X\beta\|_2^2}_{\text{Loss}} + \lambda \underbrace{\|\beta\|_2^2}_{\text{Penalty}}\end{aligned}$$

Here $\lambda \geq 0$ is a **tuning parameter**, which controls the strength of the penalty term. Note that:

- ▶ When $\lambda = 0$,
- ▶ When $\lambda = \infty$,
- ▶ For λ in between, we are balancing two ideas: fitting a linear model of y on X , and shrinking the coefficients

Example: visual representation of ridge coefficients

A visual representation of the **ridge regression** coefficients for the same example ($n = 50$, $p = 30$, and $\sigma^2 = 1$; 10 large true coefficients, 20 small) at $\lambda = 25$:



We see that squishing the coefficients toward zero gives:


- ▶ A decrease in variance!
- ▶ A decrease in MSE!
- ▶ An increase in bias...

However, ridge regression:

- ▶ Seems quite harsh on large, important coefficients.
- ▶ Leaves us with a large linear model, which is not that interpretable when we have many variables.

Ridge regression coefficients

Ridge regression coefficients, plotted against λ :



The red paths correspond to the true nonzero coefficients; the gray paths correspond to true zeros. The vertical dashed line at $\lambda = 15$ marks the point above which ridge regression's MSE starts losing to that of linear regression

An important thing to notice is that the gray coefficient paths are not **exactly zero**; they are shrunk, but still nonzero

The lasso



The **lasso**¹ estimate is defined as

$$\begin{aligned}\hat{\beta}^{\text{lasso}} &= \underset{\beta \in \mathbb{R}^p}{\operatorname{argmin}} \|y - X\beta\|_2^2 + \lambda \sum_{j=1}^p |\beta_j| \\ &= \underset{\beta \in \mathbb{R}^p}{\operatorname{argmin}} \underbrace{\|y - X\beta\|_2^2}_{\text{Loss}} + \lambda \underbrace{\|\beta\|_1}_{\text{Penalty}}\end{aligned}$$

The squared ℓ_2 penalty $\|\beta\|_2^2$ of ridge regression, has been replaced by an ℓ_1 penalty $\|\beta\|_1$. Even though these problems look similar, their solutions behave very differently

Note the name “lasso” is actually an acronym for: Least Absolute Selection and Shrinkage Operator

¹Tibshirani (1996), “Regression Shrinkage and Selection via the Lasso”



$$\hat{\beta}^{\text{lasso}} = \underset{\beta \in \mathbb{R}^p}{\operatorname{argmin}} \|y - X\beta\|_2^2 + \lambda \|\beta\|_1$$

The **tuning parameter** λ controls the strength of the penalty.


- ▶ If $\lambda = 0$:
- ▶ As $\lambda \rightarrow \infty$:

For λ in between these two extremes, we are balancing two ideas: fitting a linear model of y on X , and shrinking the coefficients. But the nature of the ℓ_1 penalty causes some coefficients to be shrunk to **zero exactly**

This leads to **variable selection**!


Example: visual representation of lasso coefficients

Our running example from last time: $n = 50$, $p = 30$, $\sigma^2 = 1$, 10 large true coefficients, 20 small. Here is a visual representation of lasso vs. ridge coefficients (with the same degrees of freedom):




Example: subset of small coefficients

Example: $n = 50$, $p = 30$; true coefficients: 10 large, 20 small




Remember this behavior for Ridge Regression?

Ridge regression coefficients, plotted against λ :




What changes?

Lasso coefficient paths



In the lasso, we get exact zeros! Moreover, the large coefficients are not impacted nearly as much!

Lasso coefficient paths



In the lasso, we get exact zeros! Moreover, the large coefficients are not impacted nearly as much!

It is often easier to view lasso results on the $\log \lambda$ scale.

Advantages of sparsity

- ▶ Interpretability: We can understand what the model relies on for prediction (understanding $\hat{f}$)
- ▶ We might gain some insight into the underlying data (though not causally) (helping to understand f)
- ▶ If we're building a predictive score, we can measure fewer things in the future (simpler $\hat{f}$ to apply later)

Sidenote: Can we do inference (e.g., p -values, confidence intervals, etc)? **Not easily!**

Important details

Just like in ridge regression:

- ▶ We don't penalize our intercept, if one is included. Again, columns are often centered and then intercept omitted.
- ▶ The penalty term $\|\beta\|_1 = \sum_{j=1}^p |\beta_j|$ is not fair if the predictor variables are **not on the same scale**. Hence, we often scale the columns of X to have variance 1.

How is it actually solved?

[REMINDER: SHOW SOFTWARE DEMOS]

Interview question

Suppose that you regress Y on X_1, X_2 using both ridge regression and the lasso (with $\lambda > 0$). If $X_1 = X_2$, what does:

The solution to ridge look like

The solution to lasso look like

Simple case: Orthogonal X

Consider the very simple case where $X = I$. We can examine how both ridge regression and the lasso behave on this data. This is easy, because the variables do not interact with one another, but it gives an intuition for their more general behavior.

When $X = I$, note that the RSS becomes

$$\|y - X\beta\|_2^2 = \|y - \beta\|_2^2 = \sum_{i=1}^n (y_i - \beta_i)^2.$$

Simple case: Orthogonal X

The corresponding penalized forms then become

Ridge:

$$\hat{\beta}_{\text{ridge}} = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^n (y_i - \beta_i)^2 + \lambda \beta_i^2$$

Lasso:

$$\hat{\beta}_{\text{lasso}} = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^n (y_i - \beta_i)^2 + \lambda |\beta_i|$$

We can minimize separately for each i !

For ridge regression, differentiating yields

and therefore

$$\hat{\beta}_i =$$

Thus the estimates from ridge regression correspond to the least squares estimates reduced by a constant multiple.

Note that this can never give zero coefficients, and that it penalizes large coefficients quite a bit.

For the lasso, it turns out that the minimizer of

$$(y_i - \beta_i)^2 + \lambda|\beta_i|$$

is given by

$$\hat{\beta}_i =$$

Thus estimates from ridge regression correspond to shrinking the least squares estimates toward zero by an additive constant (without crossing zero).

This can give zero coefficients, and is also relatively harsher on small coefficients than larger ones. It appears to be better suited to sparse settings.

Regularized regression

Least squares solves the problem:

$$\operatorname{argmin}_{\beta} \|y - X\beta\|_2^2$$

In regularized regression, we add a penalty to this squared error loss term

$$\operatorname{argmin}_{\beta} \|y - X\beta\|_2^2 + P_{\lambda}(\beta)$$

You've seen several examples: $P_{\lambda}(\beta) = \lambda \sum_j \beta_j^2$ (ridge regression) and $P_{\lambda}(\beta) = \lambda \sum_j |\beta_j|$ (lasso).

These penalties both “pull” β toward 0, causing it to be less variable. However, by making β smaller than the least squares estimate, they add bias.

The λ parameter determines how strong this penalty is. A larger λ leads to a higher bias and a smaller variance.

Regularized regression

$$\operatorname{argmin}_{\beta} \|y - X\beta\|_2^2 + P_{\lambda}(\beta)$$

$P_{\lambda}(\beta) = \lambda \sum_j \beta_j^2$ (ridge regression) and $P_{\lambda}(\beta) = \lambda \sum_j |\beta_j|$ (lasso).

These penalties both “pull” β toward 0, causing it to be less

However, by making β smaller than the least squares estimate, they add

The λ parameter determines how strong this penalty is. A larger λ leads to a higher bias and a smaller variance

Extensions of the Lasso

Adding an L_1 penalty can induce sparsity in a variety of settings.

An entire industry has been developed around producing related methods and related theory.

We will see at least:

- ▶ Logistic regression with a lasso penalty
- ▶ The elastic net

Extensions of the Lasso

Adding an L_1 penalty can induce sparsity in a variety of settings.

An entire industry has been developed around producing related methods and related theory.

Other classic examples:

- ▶ The group lasso: setting groups of variables to zero
- ▶ Graphical lasso: estimates networks of dependence between variables
- ▶ Fused lasso: estimates piecewise-constant functions
- ▶ Trend filtering: estimates piecewise-polynomial functions
- ▶ Low-rank matrix decomposition: estimate a low-rank approximation to a matrix
- ▶ Sparse additive models
- ▶ Vector auto-regressive models

Model selection and generalization error

Regularized regression is one of the first settings you see where you must pick a particular model by setting a tuning parameter (λ here).

We know that λ trades off between bias and variance, and that we want to pick this tradeoff to minimize the *generalization error* for *new* samples.

How do we actually pick this λ ?

[WARNING: We are going to talk about this for now in the i.i.d. setting. We need to be more careful for timeseries. We will come back to this.]

K -fold Cross-validation

We want to estimate how well a method with each λ should do on *new* data points.

We don't have new data points, but we can hold out a few of our own.

We split the training pairs into K parts or “folds” (commonly $K = 5$ or $K = 10$)

1	2	3	4	5
Train	Train	Validation	Train	Train

K -fold cross validation considers training on all but the k th part, and then validating on the k th part, iterating over $k = 1, \dots, K$

K -fold Cross-validation

We want to estimate how well a method with each λ should do on *new* data points.


We don't have new data points, but we can hold out a few of our own.

K -fold Cross Validation:

1. Randomly split the data into K equal-sized parts ("folds")
2. Each part is used as the test set once, with the other $K - 1$ parts combined into a training set
3. The error on all these test sets is averaged

These errors are used to select a final model (λ value) to use. The model is refit on the entire data set for this λ value.

Usually $K = 5, 10$, or n .



Standard errors for cross-validation

One nice thing about K -fold cross-validation (for a small $K \ll n$, e.g., $K = 5$) is that we can estimate the standard deviation of $\text{CV}(\theta)$, at each $\theta \in \{\theta_1, \dots, \theta_m\}$

First, we just average the validation errors in each fold:

$$\text{CV}_k(\theta) = \frac{1}{n_k} e_k(\theta) = \frac{1}{n_k} \sum_{i \in F_k} (y_i - \hat{f}_{\theta}^{-k}(x_i))^2$$

where n_k is the number of points in the k th fold


We then compute the sample standard deviation of $\text{CV}_1(\theta), \dots, \text{CV}_K(\theta)$,

$$\text{SD}(\theta) = \sqrt{\text{var}(\text{CV}_1(\theta), \dots, \text{CV}_K(\theta))}$$

Finally we estimate the standard deviation of $\text{CV}(\theta)$ —called the **standard error** of $\text{CV}(\theta)$ —by $\text{SE}(\theta) = \text{SD}(\theta) / \sqrt{K}$

Example: choosing λ for the lasso (continued)

The cross-validation error curve from our lasso example, with $\pm$ standard errors:



The one standard error rule

The **one standard error rule** is an alternative rule for choosing the value of the tuning parameter, as opposed to minimizing the cross-validation error.


We find the usual minimizer $\hat{\theta}$, and then move θ in the direction of **increasing regularization** as much as we can, while keeping the CV error within one S.E. of $CV(\hat{\theta})$:

$$CV(\theta) \leq CV(\hat{\theta}) + SE(\hat{\theta})$$

Idea: “All else equal (up to one standard error), go for the simpler (more regularized) model”

Example: choosing λ for the lasso (continued)

Back to our lasso example: the one standard error rule chooses a larger value of λ (more regularization):



Prediction versus interpretation

Remember: cross-validation error estimates prediction error at any fixed value of the tuning parameter, and so by using it, we are implicitly assuming that achieving minimal prediction error is our goal.

However, this will tend to over-select variables, complicating our **interpretation**. The one standard error rule is a step in the direction of interpretability (and a hack).

What to do next?

What do we do next, after having used cross-validation to choose a value of the tuning parameter $\hat{\theta}$?

It may be an obvious point, but worth being clear: we now fit our estimator to the **entire training set** (x_i, y_i) , $i = 1, \dots, n$, using the tuning parameter value $\hat{\theta}$

E.g., in the last lasso example, we resolve the lasso problem

$$\hat{\beta}^{\text{lasso}} = \underset{\beta \in \mathbb{R}^p}{\operatorname{argmin}} \|y - X\beta\|_2^2 + \lambda \|\beta\|_1$$


on all of the training data, with $\hat{\lambda} = 3.458$

We can then use this estimator $\hat{\beta}^{\text{lasso}}$ to make future predictions

Model assessment

Instead of just using the CV error curves to pick a value of the tuning parameter, suppose that we wanted to use cross-validation to actually estimate the **prediction error** values themselves

I.e., suppose that we actually cared about the values of the CV error curve, rather than just the location of its minimum



This is called **model assessment** (as opposed to **model selection**)

Model assessment

Why would we want to do this? Multiple reasons:

- ▶ Report estimate for prediction error to collaborator
- ▶ Compare estimated prediction error to known prediction errors of other estimators
- ▶ Answering the question: is my estimator $\hat{f}$ working at all?
- ▶ Choosing between different estimators (each possibly having their own tuning parameters)

Choice of K

The larger the choice of K , the more iterations needed, and so the more computation.

Aside from computation, the choice of K affects the **quality of** our cross-validation error **estimates** for model assessment. Consider:

- ▶ $K = 2$: split-sample cross-validation. Our CV error estimates are going to be **biased upwards**, because we are only training on half the data each time
- ▶ $K = n$: leave-one-out cross-validation. Our CV estimates

$$\text{CV}(\theta) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}_{\theta}^{-i}(x_i))^2$$

are going to be heavily positively correlated, which can lead to high variance.

Choosing $K = 5$ or $K = 10$ seems to generally be a good tradeoff

- ▶ In each iteration we train on a fraction of (about) $(K - 1)/K$ the total training set, so this **reduces the bias**
- ▶ There is less overlap between the training sets across iterations, so the terms in

$$\text{CV}(\theta) = \frac{1}{n} \sum_{k=1}^K \sum_{i \in F_k} (y_i - \hat{f}_{\theta}^{-k}(x_i))^2$$

are not as correlated, and the error estimate has a **smaller variance**

Leave-one-out shortcut

Recall that leave-one-out cross-validation is given by

$$\text{CV}(\theta) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}_{\theta}^{-i}(x_i))^2$$

where $\hat{f}_{\theta}^{-i}$ is the estimator fit to all but the i th training pair (x_i, y_i)

Suppose that our tuning parameter is $\theta = \lambda$ and $\hat{f}_{\lambda}$ is the ridge regression estimator

$$\hat{f}_{\lambda}(x_i) = x_i^T \hat{\beta}^{\text{ridge}} = x_i^T (X^T X + \lambda I)^{-1} X^T y$$

Then it turns out that

$$\frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}_{\lambda}^{-i}(x_i))^2 = \frac{1}{n} \sum_{i=1}^n \left[\frac{y_i - \hat{f}_{\lambda}(x_i)}{1 - S_{ii}} \right]^2$$

where $S = X(X^T X + \lambda I)^{-1} X^T$. This realization provides a **huge computational savings!**

A final warning about lasso

The lasso feels like magic, especially after you've tried to use regression in high-dimensions without it.

However, the lasso still selects a linear model. You will never do better than the best linear model using your variables could do for your problem.

The lasso cannot capture non-linearity or interactions (automatically). It cannot tell if your variables are poorly transformed.

Other methods can handle these problems. Do not neglect them.

What next?

We started out by noting that we want to approximate $\mathbb{E}(Y|X = x)$.

We thought about linear approximations, $\mathbb{E}(Y|X = x) = x^T \beta$, and showed that we can beat the least squares estimate with regularization to

- ▶ Reduce the variance (at the cost of bias).
- ▶ Improve interpretability by adding sparsity

However, at the end of the day, we're still using something linear!

How do we move to something with lower bias instead of lower variance?