

46-927, Spring 2021: Homework #3

Due 1:30PM EST, Wednesday, February 17, 2021.

Homework should be submitted electronically on Gradescope before the deadline. Please submit your homework as a single PDF file. We recommend using a Jupyter notebook to prepare your homework.

Note: For this homework, do all computational problems in Python. For problem 1, you can write LaTeX math in your Jupyter notebook, do a derivation in $\text{\LaTeX}$ and merge the PDF with your Jupyter PDF, or make a **legible** scan of handwritten work and add it to your PDF.

Please post any questions on the Piazza discussion board.

1. *Different forms of SVM.* An important piece of optimization is transforming one optimization problem to another equivalent optimization problem. We will consider optimization problems to be equivalent if they give the same solution at their optimum point, or if you can easily transform between their solutions.

We're going to go through a series of steps transforming the SVM optimization into other equivalent forms. Our end goal will be to see that the SVM can equivalently be rewritten as the minimization of

$$\sum_{i=1}^n (1 - y_i(x_i^T \beta + \beta_0))_+ + \frac{\lambda}{2} \|\beta\|_2^2$$

where $(z)_+ = \max(z, 0)$, the positive part function. This is the most classic form of the SVM, and has an interesting interpretation which we will discuss next class.

Note: This is a very important and classic style of approach/proof. I'm going to work through much of it with you, since you haven't seen these before. Please read everything carefully anyway and give a lot of attention to what's happening, even though you're not doing it from scratch. I've **Bolded** questions that you need to answer as you go along.

You will get a chance to check your Form 2–5 of the problem when you reach the end, since Form 5 should exactly match the equation above. You can go back and tweak your choices at the other steps so that everything works out.

We start with

Maximize M (Form 1)

Subject to $\sum_{i=1}^n \varepsilon_i \leq C$, $\varepsilon_i \geq 0$, $y_i(x_i^T \beta + \beta_0) \geq M(1 - \varepsilon_i)$, $\|\beta\|_2 = 1$

Note that this is exactly how we defined the SVM in class (with an explicit intercept).

- (a) Consider the inequality $y_i(x_i^T \beta + \beta_0) \geq M(1 - \varepsilon_i)$. Since we know $\|\beta\|_2 = 1$, we can rewrite this as

$$y_i(x_i^T \beta + \beta_0) \geq M(1 - \varepsilon_i)\|\beta\|_2$$

without changing anything.

Show that if you multiply β_0 and β by a fixed constant b , nothing changes about the meaning of this inequality.

Since you only really care about the direction of β , this means that you can set the length of β to be anything you want without changing the rest of the optimization problem.

Choose a value of $\|\beta\|_2$ (the Euclidean length of β) that lets you eliminate M from your optimization problem. **Rewrite** Form 1 using this substitution, so that M no longer appears in the problem. (Note, you can switch the “Maximize” to a “Minimize”, and you can drop the $\|\beta\|_2 = 1$ constraint). *Call your new optimization “Form 2.”*

Question: Your M parameter used to carry information about the width of your margin. Where is this information now encoded?

Question: What is the relationship between the β that solves this problem and the β that solved the previous optimization problem?

- (b) Now we’re going to work with the $\sum_{i=1}^n \varepsilon_i \leq C$ term. Suppose that we added a term $D \sum_{i=1}^n \varepsilon_i$ to the objective from Form 2, so that your minimization objective looked like

$$(\text{Form 2 objective}) + D \sum_{i=1}^n \varepsilon_i$$

and remove the constraint that $\sum_{i=1}^n \varepsilon_i < C$.

Write out the new formulation where this substitution has been made, and call it “Form 3”.

Argue that for any C in Form 2, there exists a D in Form 3 that gives the same solution to the optimization problem. (Just a sentence or two.)

Question: If I make C bigger in Form 2, does D get bigger or smaller in the equivalent problem of Form 3?

- (c) We’re getting close! Now comes the cool and tricky part. Think about a set of parameters that solves Form 3. In particular, think about a particular i and ε_i . The “Subject to” piece of Form 3 can be thought of as two constraints on ε_i , if everything else is held fixed.

Rewrite these as two inequality constraints on ε_i , where ε_i appears alone on one side.

Now think about the objective function. At the optimum, you know that ε_i should be as small as possible, subject to the constraints, because of the $\sum_{i=1}^n \varepsilon_i$ term in the objective.

From these two pieces, you can figure out exactly what ε_i must be at the optimum!

Write out ε_i as a function of $y_i(x_i^T \beta + \beta_0)$. You will probably want to use the “positive part” function, $(z)_+ = \begin{cases} x & x > 0 \\ 0 & x < 0 \end{cases}$.

Substitute this quantity for ε_i in your objective and write down the new optimization form, called Form 4. After you’ve made this substitution, you should no longer have any ε_i .

- (d) Finally, you should be able to divide through by a constant to get an optimization problem exactly of the form we set out to show. Call this Form 5.

Question: What is λ as a function of D to make the form identical to formula stated at the beginning of the problem?

Question: We know that, when C goes up, the margin width increases and the separating plane becomes potentially more biased but less variable. When C increases, do D (Form 4) and λ (Form 5) each increase or decrease? What happens to the length of β ?

2. *Experimenting with class weights.* We talked a couple classes ago about unbalanced data sets and issues that they can cause for classifiers. We'll use the SVM as a demonstration of one of these issues and a potential fix.

The file `generate_imbalance.py` contains a function to generate data in this problem. We'll look at a simple classification scenario under different proportions of 0 and 1s.

Hint: To get `plot_decision_regions` to work when you don't have any `y` values, you might need to assign `y` to be a constant vector of zeros or ones.

- (a) Call `gen_data(500,100)` to draw *training* set with 500 observations in class 0 and 100 observations in class 1. Call it again with `gen_data(500,500)` to generate a *test* set (with different class balances).
Fit a linear SVM on the training set. (We will use $C = 1$ throughout this problem, rather than worry about tuning.)
Make a plot of the decision boundary on the *training* set, using `plot_decision_regions` from `mlxtend.plotting`. It should look pretty reasonable.
- (b) Now make predictions on the *test* set using the SVM you trained in the previous part. Report a confusion matrix and a misclassification error. Do you notice anything strange? (Hint: Keep in mind that the test set is a balanced classification problem: there are an equal number of 0 and 1 in the data set.)
Make a plot of the decision boundary for the same classifier using `plot_decision_regions`, but this time plot it for the *test* set points. Can you see why you're getting weird classifications in this balanced problem?
- (c) Retrain your SVM on the *training* sample, but this time set `class_weight='balanced'` inside your call to `SVC()`. (Read the documentation for `SVC` to see how this option is explained.) This puts greater losses on points from your rare class to counteract your imbalance.
With this new classifier, reproduce:
- Your confusion matrix on the test data.
 - Your misclassification error on the test data.
 - Your plot of the decision boundary on the test data.

How have all of these changed?

3. *Anomaly detection with SVMs.* When we talked about Mahalanobis distance, we hinted at its use for detecting rare events. It turns out that the SVM can also be modified for this purpose. For this assignment we'll explore the idea.

A one-class SVM classifies whether points are “like your data” (guesses 1) or not (guesses 0). It is trained on a bunch of samples X , but no labels, since every observation is essentially a 1. Otherwise everything works like the SVM you're familiar

with. The other difference is that the cost tuning parameter `C` is replaced by a similar tuning parameter `nu`. You'll explore `nu` a bit in this problem as well.

To fit the SVM, use the function `OneClassSVM` from `svm` in `sklearn`. For example:

```
svm.OneClassSVM(nu=0.1, kernel="rbf", gamma=0.5).fit(X,y)
```

The file `anomaly.py` contains starter code to generate data for this problem. You're going to start by trying a few different settings to see how this method behaves. Generate a data set with `get_data(500)`.

- (a) Create a scatter plot of your data. It should not look anything like a Gaussian, so Mahalanobis distance would not work well.
- (b) Fit the one-class SVM to your data. Use the `rbf` kernel, with `nu=0.1` and `gamma=0.5`. Plot the resulting decision regions, along with your data, using `plot_decision_regions`.
You should see a boundary around the sampled points. This is what the one-class SVM does, it tries to estimate a boundary around a small region that will tend to contain your future points.
- (c) Experiment with changing `gamma` to be higher (to 2) and lower (to 0.1) and plot the results. What changes about the resulting boundary? Give a brief (one or two sentences) description of how `gamma` can be interpreted here.
- (d) Experiment with changing `nu` to be higher (0.4) and lower (0.01), while keeping `gamma` fixed at 0.5. Plot the results. What changes about the resulting boundary? Give a brief (one or two sentences) description of how `nu` can be interpreted here.
- (e) Finally, let's actually use this! For a simple illustration, we'll use the adjusted closing prices of the DJIA from 2003–2009. We will try to look for anomalous patterns in this time series.

One simple way to look for anomalous patterns in time series is to slide a fixed window of length K along the time series, and for each position in the time series, create a length K observation in a data set. This results in a $(n - K + 1) \times K$ data set, where each observation corresponds to a chunk of time. We can then compare to this data set to see if future observations differ significantly from past chunks of K days.

To make things easier, I've provided you with `djia_transformed.csv`, where this process has already been carried out. For this example, I used $K = 4$.

Fit a one-class SVM to the first 467 rows. These correspond to 2003–2004. Use an rbf kernel with `gamma=0.0000001` and `nu=0.01` (though the results are pretty insensitive).

Now we'll make predictions on the entire data set (2003–2009). We'll try to see if anything strange happened in the market during this time...

Plot the predictions as a function of time. Because nothing is really stationary here, points will pretty quickly look like anomalies outside of your training region, making this plot a little uninteresting.

To get a better signal for anomalies, we can plot the distance from the SVM boundary (including sign) as a function of time. To get this distance, use the `decision_function(X)` method of your fitted SVM instead of the `predict` function. Plot this quantity as a function of time. A negative score is a signal that a point looks anomalous. Do you see anything particularly anomalous, and does it correspond to anything that you expect?